THE UNIVERSITY OF TULSA

THE GRADUATE SCHOOL

HYBRID ATTACK GRAPHS FOR MODELING CYBER PHYSICAL SYSTEMS

SECURITY

by
George Robert Louthan IV

A thesis submitted in partial fulfillment of

the requirements for the degree of Master of Science

in the Discipline of Computer Science

The Graduate School

The University of Tulsa

2011

THE UNIVERSITY OF TULSA

THE GRADUATE SCHOOL

HYBRID ATTACK GRAPHS FOR MODELING CYBER PHYSICAL SYSTEMS

SECURITY

by
George Robert Louthan IV

A THESIS

APPROVED FOR THE DISCIPLINE OF

COMPUTER SCIENCE

By Thesis Committee

_____, Chairperson
John C. Hale

_____
Mauricio Papa?

_____
Peter Hawrylak?

ABSTRACT

George Robert Louthan IV (Master of Science in Computer Science)

Hybrid Attack Graphs For Modeling Cyber Physical Systems Security

Directed by John C. Hale

42 pp., Chapter 1: Conclusions

(75 words)

As computer systems' interactions with the physical world become more pervasive, largely in safety critical domains, the need for tools to model and study the security of these so-called cyber physical systems is growing. This thesis presents extensions to the attack graph modeling framework to permit the modeling of continuous, in addition to discrete, system elements and their interactions, to provide a comprehensive formal modeling framework for describing cyber physical systems and their security properties.

ACKNOWLEDGEMENTS

TABLE OF CONTENTS

Page

LIST OF TABLES

LIST OF FIGURES

CHAPTER 1

INTRODUCTION

1.1   Introduction

As computer systems become pervasive across a variety of domains, not only are their interactions with people becoming more frequent; computer systems are also increasingly interacting with the physical world and with each other.

Systems that include both continuous and discrete components are termed *hybrid systems*. When linked together with a significant network component, these systems are sometimes called *cyber physical systems*. They have been targeted as a key area of research by the National Science Foundation because they are becoming pervasive in safety-critical domains such as medical, critical infrastructure, and automotive equipment. This thesis is concerned with modeling the security of these systems and their interactions with each other and the physical world.

1.2   Modeling Frameworks

An excellent argument for the need for new research in modeling cyber physical systems is due to Lee in a 2006 position paper in the National Science Foundation Workshop on Cyber-Physical Systems, a prelude to the NSF's research initiative on cyber physical systems:

Cyber-Physical Systems (CPS) are integrations of computation with physical processes. Embedded computers and networks monitor and control the physical processes, usually with feedback loops where physical processes affect compu- tations and vice versa. In the physical world, the

passage of time is inexorable and concurrency is intrinsic. Neither of these properties is present in today's computing and networking abstractions. [17]

Existing frameworks for modeling and analysis of purely discrete computer networks are inappropriate for use in these systems because of their inability to capture the continuous domain; they also lack a robust, let alone "inexorable" notion of time. Likewise, existing methods for studying hybrid systems fall short when it comes to modeling the complex distributed networks that are often the hallmarks of cyber physical systems.

1.3   Scope

This thesis presents an extension of the attack graph modeling framework, a discrete domain formalism for studying network security, into the continuous domain to enable it to model cyber physical systems. The goal is to combine aspects of both hybrid systems modeling frameworks, particularly hybrid automata, which best describe systems in relative isolation; and computer network security modeling frameworks, particularly attack graphs, which excel at capturing the complex interrelationships and interdependencies among assets and attacks.

The remainder of this thesis is structured as follows. Chapter 2 provides background in hybrid systems and their modeling methods, introduces past work in attack graphs, and presents a set of case studies in both the hybrid and discrete domains to be used throughout this work. Chapter 3 introduces in detail the attack graph framework to be used as the basis for the hybrid extensions. Chapter 4 introduces the extensions themselves. Chapter 5 delivers some results from this modeling methodology, and Chapter 6 draws conclusions and suggests further work.

CHAPTER 2

BACKGROUND

2.1 Cyber Physical Systems

2.1.1 Hybrid Systems

A system with both continuous (frequently physical) components and discrete (frequently digital) components is said to be a *hybrid system*, named for its characteristic blending of the two domains. Examples of hybrid computer systems abound in industrial controls, for example, although hybrid systems may also be fully physical (e.g., a bouncing ball that experiences continuous behavior when rising and falling and discrete behavior when colliding with a surface).

The term hybrid system is an older one that was coined as researchers began to study the newly pervasive reactive systems that arose as programmed control of the physical world became widespread [3]. For several reasons it does not suffice to describe precisely the types of systems with which this work is concerned: a subset of hybrid systems that incorporate a significant computer and networking component.

Nevertheless, the modeling of hybrid systems is well studied and provides a sufficient body of relevent work from which to draw to warrant its inclusion. This chapter includes background on a particularly relevent modeling framework for hybrid systems called the hybrid automaton, which is used in this thesis as the standard benchmark against which to compare hybrid modeling techniques.

2.1.2 Definition

A newer, better term for the systems investigated in this thesis is *cyber*

physical systems. Put simply, a cyber physical system is a networked hybrid system: a networked computer system that is tightly coupled to the physical world.

2.1.3 Challenges

According to the 2008 Report of the Cyber-Physical Systems Summit, "The principal barrier to developing CPS is the lack of a theory that comprehends cyber and physical resources in a single unified framework." [1]

The summit further identified as part of the necessary scientific and technological foundations of cyber physical systems both (1) new modeling frameworks that "explicitly address new observables" and (2) studies of privacy, trust, and security including "theories of cyber-physical inter-dependence" [1], a major theme of this work.

Crenshaw and Beyer recently enumerated four principal challenges in cyber physical systems testing that are equally apt for security: their concentration in safety critical domains, their frequent integration of third-party or otherwise unrelated systems, their dependence upon unreliable data collection, and their pervasiveness [9].

2.1.4 Hybrid Automata

Definition: A valuable formalism for modeling hybrid systems in isolation and with limited composition is the hybrid automaton of Alur, et al. [3]. This section introduces the version of the formalism described in 1996 by Henzinger [12], to which a reader interested in more than a superficial understanding is referred.

Formally, a hybrid automaton H is made up of a set of real-valued state variables, their first derivatives, a set of operational modes and switches between the modes, and predicates attached to those modes and switches describing the operation of the system in those modes and the discrete transitions between them. One can think of a hybrid automaton as a pairing of a finite state machine whose
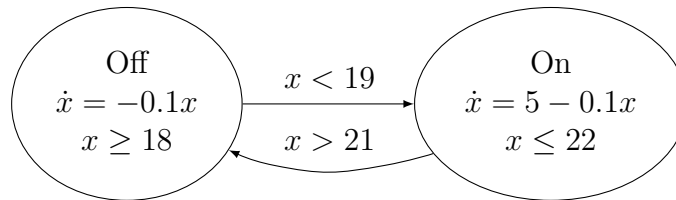
Figure 2.1: Thermostat hybrid automaton

states (called modes) and transitions (called switches) denote the discrete-domain behavior of the hybrid system, with a set of differential equations attached to each mode, which govern its continuous-domain behavior. Switches may also be labeled in order to permit synchronization across composed hybrid automata.

Modes may be decorated with invariant conditions (which state whether the system is allowed to be in that mode), flow conditions (which state how the continuous domain state variables are permitted to evolve while in that mode), and initial conditions (which state under which, if any, conditions the automaton may begin its operation with that mode). Switches are decorated with jump conditions, which serve as guards on the switch determining both (1) when the switch is allowed to be taken, and (2) the discrete changes in state variables due to that switch's activation.

A simple example of a hybrid automaton is given in Fig. 2.1, which models a simple heater thermostat [12]. The vertices in the automaton represent its operating modes, and the edges represent its control switches. In the "Off" mode, the temperature (given by x) must be greater than or equal to 18, and its first derivative with respect to time (denoted $\dot{x}$) is $-0.1x$, which represents a cooling of the environment. When the temperature is strictly less than 19, the switch from off to on is available (but not mandatory until the off mode's invariant condition $x \geq 18$ ceases to be satisfied.) The switch from on to off behaves similarly.

The hybrid automaton model is sufficiently rich to capture many hybrid systems.

Shortcomings: There are some problems with the hybrid automaton model. A hybrid automaton is not guaranteed to have a valid execution, and computing whether it does or not is non-trivial [20]. Model checking has been developed for only some subclasses of automata [13] [11], and many desirable properties of them are undecidable [14].

However, there are even more nagging problems when considering hybrid automata or their variants for the study of cyber physical systems. One of the hallmarks of cyber physical systems is a distributed and highly networked nature. While they provide a natural model for the discrete-continuous boundary, hybrid automata have only a rudimentary notion of communication, no clear means for specifying message passing, and when used in large topologies have significant scaling problems, both computationally and cognitively.

Alternatives: Some attempts have been made to solve the problem of the hybrid automaton's unsatisfactory capability for modeling networks and communication. Particularly, the designation of shared actions and shared variables as "input" or "output" is a popular tactic, used in the powerful hybrid I/O automaton [22] [21], its descendent the timed I/O automaton [16], and also in the PHAVer model checker [11].

The work of this thesis is also something of an outgrowth from an instance of this strategy in which prototypical "hybrid link automata" were developed to model explicit communication channels. An example of the cognitive scalability issues inherent with this design is given in Fig. 2.2, a considered "hybrid link automaton" prototype modeling a link on which messages may be dropped, injected, or delayed, and on which rudimentary mutual exclusion of messages is enforced. This strategy may have a place in modeling some systems but falls short of the goal of modeling complex, interdependent networks of hybrid systems with more
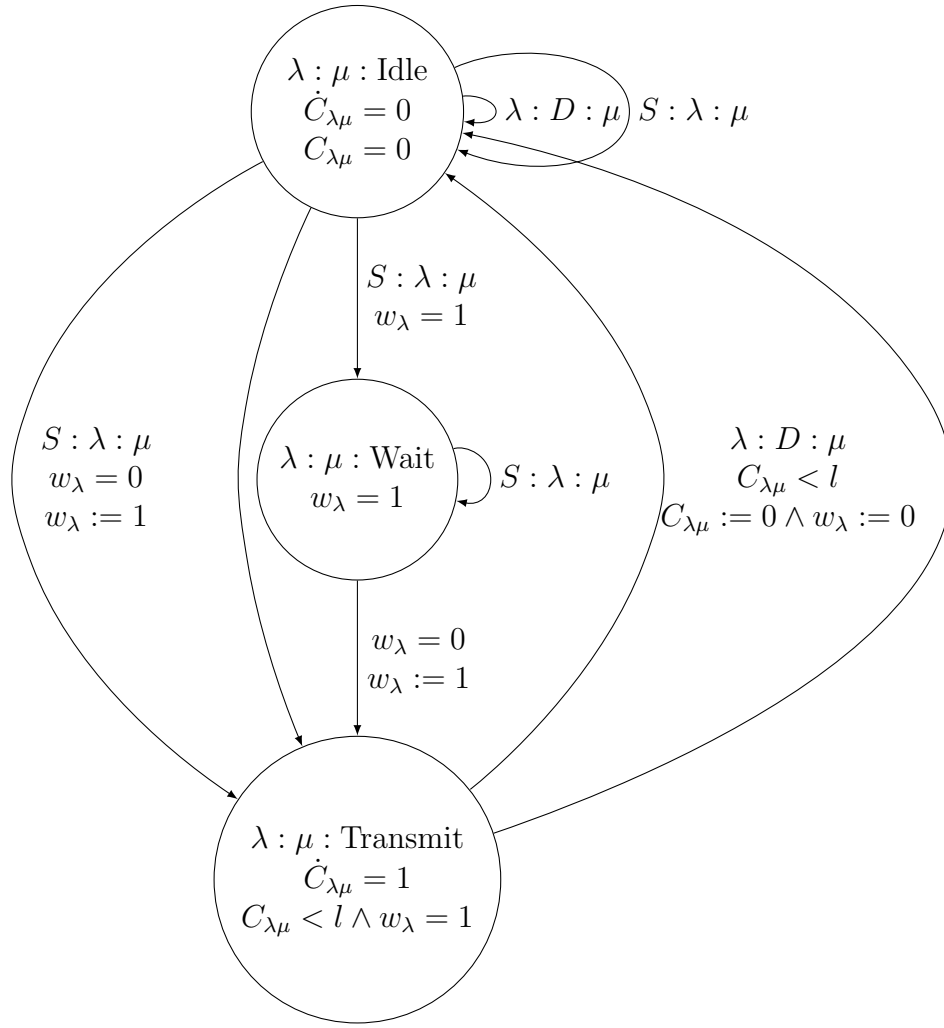
Figure 2.2: Example hybrid link automaton

conventional computer networks.

2.2 Attack Graphs

2.2.1 Introduction

An attack graph is one of several related formalisms that utilize graph theory to model the state space of computer systems attacks. Perhaps they are best introduced when presented as an alternative to a similar model called an attack tree.
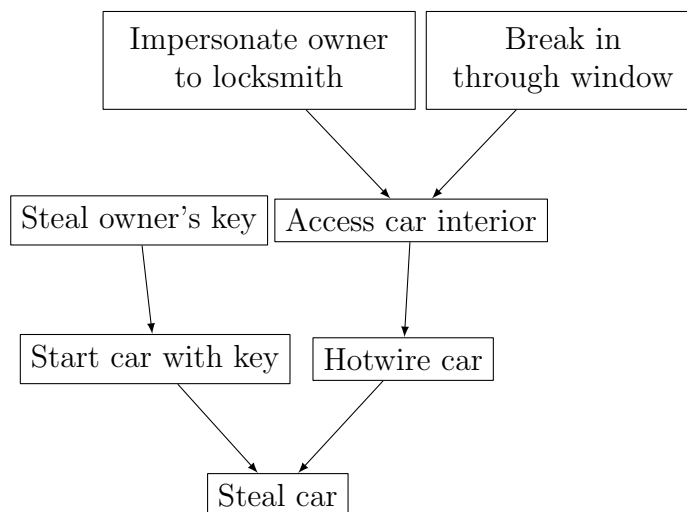
Figure 2.3: Simple car theft attack tree

2.2.2 Attack Trees

An attack tree is a goal-oriented tree model of an abuse of a system [25]. The root of the tree represents the attacker's goal, and the children of any given node represent the prerequisite activities required to reach that node. For example, consider the goal of stealing a car, which is modeled in a simple attack tree in Fig. 2.3.

The attacker must start the car and drive away; this could be accomplished either by breaking in and hotwiring the car, or by stealing the owner's key and using it to subsequently steal the car. The root of the tree represents the final goal of the theft, with prerequisite goals flowing upward from the leaf nodes.

There are a few features of this modeling method to note. It is goal oriented, meaning that the consequences of the attack are known, and the goal is to ennumerate and and analyze the means by which those consequences could be reached. It is, as an attack model, agnostic to the underlying system model which makes it difficult to generate automatically. Finally, and perhaps most significantly, it captures the ways in which an attacker's actions interact and depend upon each other.

This threat-centric model is not necessarily the most useful for system stake-

holders. It requires, in a sense, that one work backward from the attack to the system state necessary to realize the attack. If, instead, an analyst desires to work from a system characterization and explore the attack space permitted by that system characterization, the attack tree framework must be in some sense turned upside down. Attack graphs do exactly that.

2.2.3   Attack Graphs

Introduction: In contrast to attack trees, attack graphs permit a topology-aware exploratory analysis of the state space of a system. It is a graph theoretic model in which vertices represent individual system states, and edges represent state transitions caused by an adversary. The concept as introduced in 1998 included notions of generalized attack patterns to be bound to state transitions; network elements and their individual configurations; network topology (three characteristics common to all current attack graph iterations); a notion of the attacker's capabilities, and edge weights representing likelihood [23]. A similar structure called a privilege graph was introduced in 1994 [10].

Most approaches to attack graph modeling represent exploits (attack patterns) as using preconditions and postconditions [18] since this was suggested in about 2000 [27]. Exploits are chained together by matching preconditions in a state node's underlying system model and applying their postconditions to generate a successor state.

Model Types: The modeling substrates of attack graphs can be broadly separated into two schools of thought, separated by the philosophy that guides the representation of the underlying network model over which network states and transitions are computer.

A specification of an underlying network model may be done with only very

9

loose restrictions, allowing arbitrary keywords as named qualities and topoligies of network objects. This thesis employs this method. It is also favored in the work of George Mason University [4] [29]. It has the advantage of permitting more straightforward adaptation into the continuous domain, which is the reason it is favored by this work.

An alternate specification method is much more restricted, confining the modeler to certain sets of terms that may, for example, impose explicit computer networking comcepts onto the model [27]. This permits generation and analysis to take a more nuanced view of a network state, including reachability analysis to determine whether a given topology permits communication between two hosts [15]. This approach is favored in the work of MIT Lincoln Laboratory and the University of California, Davis.

Generation: Attack graph generation is the process of chaining exploits to enumerate the attack space [7] [23] [26]. Methods for generating attack graphs share a common general architecture among the modern methods that use preconditions and postconditions in exploit definitions, pictured in Fig. 2.4. The attack graph generation process combines network state and exploit patterns as input, applying exploit postconditions back onto the network state to generate its output of successor states.

Research Directions: Research in attack graphs is spread throughout a variety of pathways. These include evaluating a network's security [4], specification of formal languages to represent attack graphs [27], intrusion detection system integration [28], automatic generation of security recommendations [29], and reachability analysis between hosts in a single network state [15]. For a thorough literature review up to 2005 and more detailed discussion of popular research directions, refer to the work of Lippmann and Ingols [18].
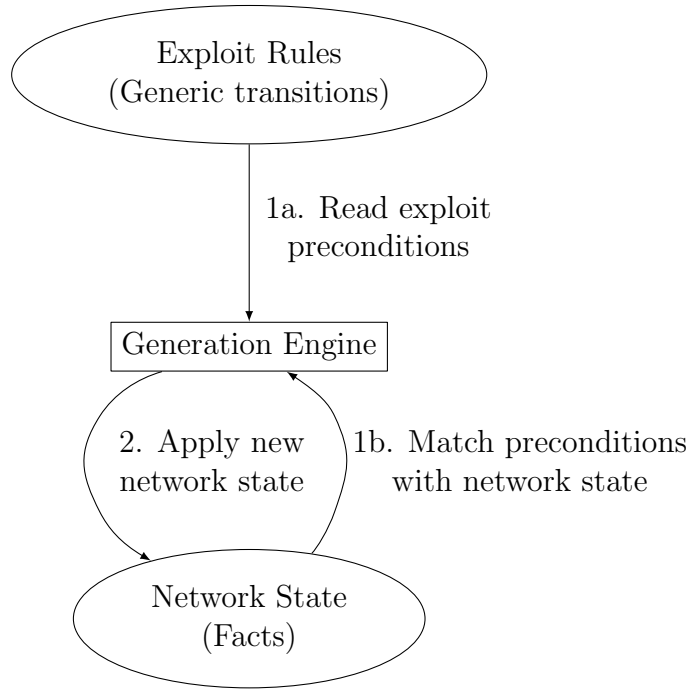
Figure 2.4: Attack graph generation process

2.3 Case Studies

Throughout this thesis, two examples of attacks are used to illustrate the models presented. The first is on a traditional information system, based upon an offensive educational exercise deployed at the University of Tulsa in 2008 involving several chained attacks. The second is the denial of service through battery exhaustion of a simple cyber-physical system of active radio frequency identification (RFID) tags and readers.

2.3.1 Blunderdome

The first case study is an attack on a simulated educational network deployed as part of a security engineering course in 2008. Dubbed the Blunderdome, it featured a firewalled network of two hosts available per attacker. See Table 2.1 for a listing of the stages and their preconditions and results. The attacker was required to log into a login server by cracking its weak SSH key (due to an operating system vulnerability), execute an elevation of privilege (due to a Linux kernel vulnerability),

Figure 2.5: Blunderdome network architecture

log into the web server, and execute a SQL injection attack to change a simulated grade. The architecture from the exercise is provided in Fig. 2.5 [19].

2.3.2 RFID Denial of Sleep

The second case study used throughout this work is a denial of service attack on the ISO 18000-7 RFID tag inventory system similar to those used by the United States Department of Defense for shipping tracking and the Department of Energy for tracking spent fuel containers [8]. The attack is similar to the ones described by Buennemeyer, *et al.* [6], and is of a newly distinguished class of attacks sometimes termed denial of sleep attacks [5] [24].

These ISO 18000-7 RFID tags are active and battery powered; they are used for inventory and shipment tracking. In particular, they are used by the Department of Energy to monitor the location and seal status of radioactive material containers, greatly reducing workers' radiation exposure. The batteries on the tags should last as long as possible in order to limit radiation exposure to maintenance workers, and the loss of power to these devices has severe safety and security consequences. An

12

Stage	Precondition	Attack	Postcondition
Gain remote user access	SSH public key available (given); weak public key	Break weak public key	User privileges on login server
Gain root access	User-level access	Execute `vmsplice` privilege escalation	Root privileges on login server; access to web server credentials
Change grade	Address and credentials for web service	Execute SQL injection	Altered grade in database

Table 2.1: Stages of the Blunderdome attack

energy draining attack to deplete the tags' batteries could significantly speed this loss of power.

The ISO 18000-7 tags have two modes: an active mode, and a sleep mode in which their power consumption is significantly reduced. The active mode has a 30 second timeout, which will cause them to sleep unless the timer is reset by the receipt of a valid command from the reader or a wake-up signal. In sleep mode, the tags will only respond to a wake-up command, which causes them to enter active mode. A denial of sleep attack occurs when the tag is not permitted to enter sleep mode or is awoken more frequently than normal.

This attack can be realized in two ways. The first is that a second, rogue RFID reader is placed by the attacker within range of some or all of the active tags. The second is for the attacker to compromise an existing reader by hacking into a computer system connected to it via a network. This presupposes that the reader is connected to the Internet or some private network into which the attacker can intrude.

Hybrid automata serve to represent the behavior of the devices themselves quite well. Fig. 2.6 represents the reader (legitimate or rogue), which does nothing

WakeUp

Command

Reading
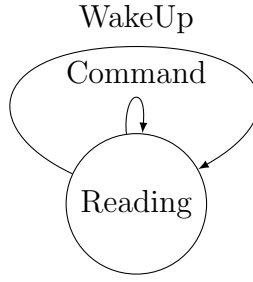
Figure 2.6: Hybrid automaton model of the RFID reader

but transmit commands and wake-up signals, which can come at any time with no restrictions. Under ordinary operating conditions this might be a few times per day over the course of several years before the batteries in the tags are drained [8].

Fig. 2.7 depicts a model of the tags. For simplicity, they are shown as starting in the active mode. It has two state variables: c, which represents the active mode timeout clock, and B represents the capacity of the battery. In active mode, the battery drains at a rate of -50 per second, a rate chosen arbitrarily for illustrative purposes only. In sleep mode, the battery drains at a rate of -1 per second. No restrictions are placed on the starting condition of the battery. The two automata are composed using two shared actions: *WakeUp* and *Command*, which synchronize the switches they decorate between the automata.

$$c = 30$$

Awake
$$\dot{c}_{\lambda\mu} = -1 \wedge \dot{B} = -50$$
$$c > 0 \wedge B > 0$$

Command
$$c := 30$$

WakeUp
$$c := 30$$

Timeout
$$c = 0$$

WakeUp
$$c := 30$$

Asleep
$$\dot{c} = 1 \wedge \dot{B} = -1$$
$$B > 0$$

Command

Die
$$B = 0$$

Die
$$B = 0$$

Dead
$$\dot{c} = 0 \wedge \dot{B} = 0$$
$$B = 0$$

Command

WakeUp
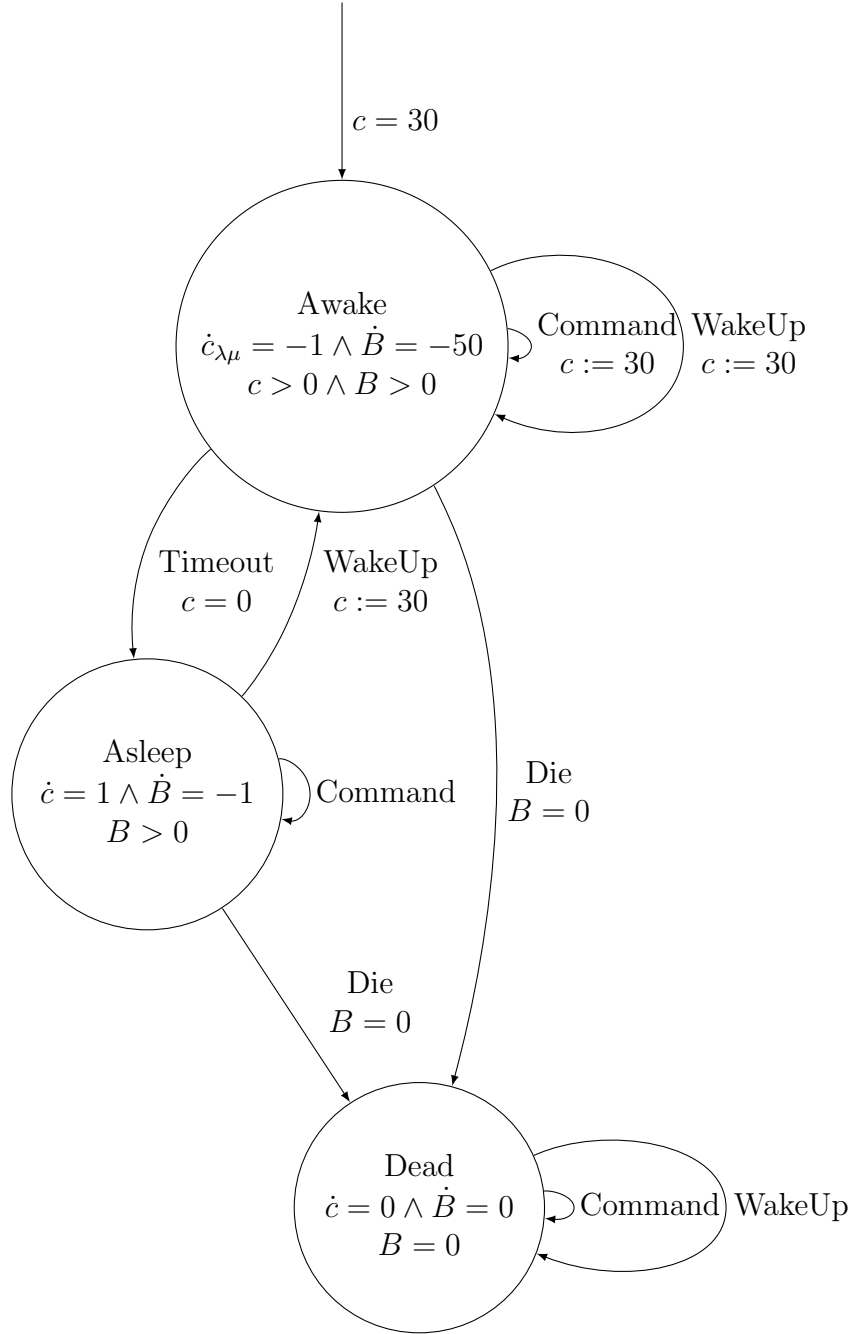
Figure 2.7: Hybrid automaton model of the case study active RFID tags

CHAPTER 3

ATTACK GRAPHS

3.1   Introduction

This chapter presents a version of the attack graph modeling framework specific to traditional information systems. The framework has a very permissive model that includes notions of assets, qualities, topologies. Assets represent potentially attackable system components; qualities assign an arbitrary string value to a named property of an asset, and topologies bind pairs of assets together with a named connection between them. Exploit patterns with preconditions and postconditions attached to free variables that can be bound at generation time to assets serve to describe potential state transitions. The next section contains a more formal definition of the basic framework.

Because of the extremely permissive model for naming properties and topologies in this scheme, the modeler must fix a lexicon of terms and their meanings in order to proceed systematically. Section 3 of this chapter presents a preliminary version of such a lexicon, fixing enough terminology to proceed with the case studies.

Research into attack graphs can be split into four broad categories:

Modeling Attack graph modeling concerns the development of the underlying representation and use of that representation to model systems. Terminology belongs here, as do efforts to automatically generate network models from real networks and exploit patterns from vulnerability databases.

Generation Attack graph generation is the process of building a graph out of a model by closing the state space over its exploits. This is where most

performance work is concentrated. A good portion of the work that enables a representation of time to be included in attack graphs belongs here as well. Constraints (such as monotonicity) on the way that state transitions are allowed to progres also fall under the generation category.

Analysis Analysis of attack graphs focuses on drawing conclusions about a system based upon the attack graph generated from its model. Work here includes integration with intrusion detection systems, automatic delivery of mitigation recommendations, and the identification of security consequences with particular states.

Visualization Visualization of attack graphs seeks to reduce or eliminate their known cognitive scalability issues; the goal is to deliver the results of the other three steps in a meaningful fashion.

This thesis is mainly concerned with the modeling stage, its goal being to introduce an architecture for modeling hybrid systems with attack graphs. However, some amount of work is included in the generation stage, particularly some thoughts on dealing with the progression of time; consideration is also given to analysis, particularly on identifying failure states and on aggregating sufficiently similar states (which also touches visualization).

Sections 4 and 5 introduce two analysis tools for attack graphs: state predicates and state aggregation. State predicates are logical predicates concerning a network state that could be used to specify particular failure cases, nominal system behavior, operating modes, or similarity conditions that permit states to be visually grouped due to some shared characteristics. This concept is necessary to allow the following chapter's introduction of time to be tractable.

3.2 Definition

17

For the purposes of this work, an attack graph is comprised of the following components.

Assets Assets are the subjects in the attack graph formalism, mainly representing attackable system components. For example, in a model of the RFID denial of sleep attack, the assets are the RFID tags, the reader, the possible rogue reader, and any hosts connected over the network with the RFID reader. Assets are specified with unique names, and they are decorated with *qualities* and *topologies*. Assets can also be used to model users and adversaries if it is necessary to give them explicit properties. A model's collection of assets is fixed at definition time and is not changed by state transitions.

Qualities Qualities represent properties of an asset, such as a software package or version that is installed, whether it is in sleep mode or not, and so on. Qualities are can be considered a key/value pair, where both the key and the value are string tokens. For example, the `host1` asset may have the quality `power` and the value `on`. Together with topologies, qualities make up the collection of facts.

Topologies Topologies represent relationships between two assets. These can be physical such as denoting that a printer is plugged into a computer, logical such as denoting that one host is accessible from another on an adjacent network, or more abstract such as a particular trust relationship or level of access that a subject has on another. Topologies are directed and named with string tokens. For example, `host1` might be accessible over the network via the web by `host2`, so a directed topology from `host2` to `host1` called `network_remote_web` might be used. Together with qualities, topologies make up the network state's collection of facts.

State A network or system state is comprised of all of the facts about the system's asset collection. A network model's state is fully described by the asset collection and the fact base; given the constant asset collection, a state is uniquely described by its fact base. The fact base is all the qualities and topologies that are valid for that state.

Exploit patterns Exploit patterns are generalized templates for how the actions of the attacker can alter the system state by inserting and removing qualities and topologies (but not assets). They are written as functions that take a number of parameters corresponding to assets, mapping a set of preconditions (facts about the free asset parameters) to a set of postconditions: insert and delete actions on qualities and topologies to update the fact base and therefore generate a new network state.

3.2.2 Formal

Primitive Domains: In order to provide a more formal description of the attack graph modeling framework, it is necessary to fix first the domains of note. The primitive domains are the most basic domains that describe the atomic units of the formalism: assets, properties (qualities), values (qualities), relationships (topologies), vulnerabilities (exploit pattern identifiers), and parameters (free asset variables in the attack patterns), and operations (used in postconditions representing insert or

delete actions):

$$\mathcal{A} : \text{assets}$$

$$\mathcal{P} : \text{properties (quality names)}$$

$$\mathcal{V} : \text{values (property values)}$$

$$\mathcal{R} : \text{relationships (topology names)}$$

$$\mathcal{W} : \text{vulnerabilities (exploit pattern names)}$$

$$\mathcal{I} : \text{parameters (free asset names)}$$

$$\text{Op} = \{\text{ins}, \text{del}\}$$

Compound Domains: Compound domains comprise the fundamental concepts that are composed of combinations of members of the primitive domains. These form the level of abstraction that it is most convenient to discuss in the articulation of the execution model and in the preceding intuitive definition, for instance.

Qualities Qualities bind an asset to a property to a value; therefore their domain is the cartesian product of those domains. The n subscripts denote that these are bound qualities of a network state, rather than free qualities in exploit preconditions and postconditions:

$$\mathcal{Q}_n : \mathcal{A} \times \mathcal{P} \times \mathcal{V}$$

Topologies Topologies bind an asset to another asset through a relationship; therefore their domain is the cartesian product of those domains:

$$\mathcal{T}_n : \mathcal{A} \times \mathcal{A} \times \mathcal{R}$$

Network states A network state, then, is denoted by a collection of assets, and

a fact base of qualities and topologies; the domain of a network state is the cartesian product of the power sets of these domains:

$$\mathcal{N} : \mathbb{P}(\mathcal{A}) \times \mathbb{P}(\mathcal{Q}_n) \times \mathbb{P}(\mathcal{T}_n)$$

Exploit patterns Exploit patterns, taken from the domain $\mathcal{E}$ depend upon free versions of qualities and topologies, which are parameterized by members of $\mathcal{I}$ rather than $\mathcal{A}$, which are used in preconditions and, when combined with an operator, postconditions:

$$\mathcal{Q}_e : \mathcal{I} \times \mathcal{P} \times \mathcal{V}$$

$$\mathcal{T}_e : \mathcal{I} \times \mathcal{I} \times \mathcal{R}$$

$$\text{Preconditions } \mathcal{P}rc_e : \mathbb{P}(\mathcal{Q}_e) \times \mathbb{P}(\mathcal{T}_e)$$

$$\text{Postconditions } \mathcal{P}oc_e : \mathbb{P}((Op, \mathcal{Q}_e)) \times \mathbb{P}((Op, \mathcal{T}_e))$$

$$\mathcal{E} : \mathcal{W} \times \vec{\mathcal{I}} \times \mathcal{P}rc \times \mathcal{P}oc$$

Attacks An exploit pattern whose parameters have been bound to assets is referred to as an attack. It takes a similar appearance:

$$\text{Preconditions } \mathcal{P}rc_n : \mathbb{P}(\mathcal{Q}_n) \times \mathbb{P}(\mathcal{T}_n)$$

$$\text{Postconditions } \mathcal{P}oc_n : \mathbb{P}((Op, \mathcal{Q}_n)) \times \mathbb{P}((Op, \mathcal{T}_n))$$

$$\mathcal{X} : \mathcal{W} \times \vec{\mathcal{I}} \times \mathcal{P}rc \times \mathcal{P}oc$$

3.3   Execution Model

3.3.1   Introduction

The reference implementation of this thesis's attack graph execution model

does not use this formal notation to represent system elements, instead favoring a more user friendly specification language. This section describes that specification language and the generation process.

3.3.2 Network Model Specification

A network model specification consists of a list of assets and an initial fact base (list of qualities and topologies). It begins with the phrase `network model`, followed by the `=` symbol.

The first component of the specification itself is an asset list. The asset list is a semicolon separated list of the names of network assets preceded by the word `assets` and a colon. For example:

```
network  model =
assets  :
asset_1 ;
asset_2 ;
asset_3 ;
```

Following the asset list, the initial fact base is specified as a semicolon separated list of facts freceded by the word `facts` and a colon. Facts may occur in any order. Qualities are specified by the word `quality`, followed by a colon, followed by the asset in question, a comma, then the name of the quality, then the `=` symbol, then the value of the quality. Topologies are specified by the word `topology`, followed by a colon, followed by the names of the assets in question separated by a directionality symbol: `->` to represent a one-way topology from the left asset to the right asset, or `<->` to represent a two-way topology. This is followed by a comma, then the name of the topology. The fact listing, and thus the network model specification, is concluded with a period. For example:

```
facts  :
```

```
         quality : asset_1 , quality_1=value_1 ;

         quality : asset_2 , quality_1=value_2 ;

         topology : asset_1 —>asset_2 , topology_1 ;

         topology : asset_2 <—>asset_3 , topology_2 ;
```
.

The completed version of this network model specification would be as follows.

```
network  model =

    assets  :

    asset_1 ;

    asset_2 ;

    asset_3 ;


    facts  :

         quality : asset_1 , quality_1=value_1 ;

         quality : asset_2 , quality_1=value_2 ;

         topology : asset_1 —>asset_2 , topology_1 ;

         topology : asset_2 <—>asset_3 , topology_2 ;
```
.

3.3.3   Exploit Specification

Exploits are specified using a related scheme. An exploit pattern resembles a function and contains four parts: a header or signature, preconditions, postconditions, and a terminating period symbol.

An exploit specification begins with the word `exploit` followed by a unique identifier to name the exploit pattern, an opening parenthesis, a comma separated list of parameter names (which are to be bound to assets), a closing parenthesis, and the `=` symbol. For example:

```
exploit  exploit_1(asset_param_1 , asset_param_2)=
```

The preconditions list follows. It begins with the word `preconditions`, followed by a colon, followed by a semicolon separated sequence of facts, which adhere to the same grammar as the facts for network model specification. For example:

```
preconditions:
    quality:asset_param_1 , quality_1=value_1;
    topology:asset_param_1 ->asset_param_2 , topology_1;
```

The postconditions list follows this. It begins with the word `postconditions`, followed by a colon, followed by a semicolon separated sequence of operations. Operations consist of the word `insert` or `delete`, followed by a space, followed by a fact. For example:

```
postconditions:
    delete  topology:asset_param_1 ->asset_param_2 , topology_1;
    insert  quality:asset_param_1 , quality_1=value_2;
```

The exploit pattern is terminated with a period character. The full example from this subsection would appear as follows.

```
exploit  exploit_1(asset_param_1 , asset_param_2)=
    preconditions:
        quality:asset_param_1 , quality_1=value_1;
        topology:asset_param_1 ->asset_param_2 , topology_1;
    postconditions:
        delete  topology:asset_param_1 ->asset_param_2 , topology_1;
        insert  quality:asset_param_1 , quality_1=value_2;
    .
```

Before continuing, a few words on the semantics of exploit processing may reduce the chance of confusion. The `insert` operation inserts a new rule in the fact base, which has the effect of replacing any previous rule with which it conflicts. For example, the insertion of `value_2` as the value of the quality `quality_1` in the example exploit `exploit_1` would replace the existing value of `value_1` specified in the preconditions; no ambiguity is introduced. The same is true of topologies.

Furthermore, the model has no innate distinction between one-way and two-way topologies except the `<->` shorthand for specifying symmetric topologies. That is to say, a bidirectional topology fact is actually implemented as two unidirectional topology facts. Therefore, for example, if a bidirectional topology exists in the fact base, one of its component directional topologies may be removed by the realization of an exploit.

3.3.4 Generation Process

This thesis is primarily concerned with modeling attack graphs, but it is worth giving some consideration to the attack graph generation process used in its reference implementation.

Currently attack graph generation proceeds according to a straightforward process depending upon a monotonicity assumption: the attacker never moves "backwards". That is, once an exploit is realized, even though the attacker may have the capability of undoing it, he does not do so. The implementation used for this thesis is written in Erlang and makes extensive use of Erlang's built-in `digraph` (its directed graph module), used to represent the attack graph itself, and `ets` (Erlang term storage, which provides constant-time table lookups) to associate graph vertices with network state data.

A maximum attack graph "depth" (really maximum permitted shortest path length from the node representing the initial state) is selected before generation

begins. Then the program initializes a graph with a single node (the starting state), and places the starting state and the node identifier in a lookup table. Then it begins to recurse using the depth value, which is initialized to the value selected by the user; and a list of available states (nodes), which is initialized to contain a single entry: the starting state.

The generation process is as follows. For each available state, a list of all valid attacks (or bound exploit patterns) is generated. Every permutation of every possible exploit realization from the state according to the pattern preconditions is included in this attack list. The attacks' postcondition operations are applied to the network state in order to generate (or load from memory, if the new fact base has already been generated by another attack on this or another state) successor states and add directed edges from the state in question to its successor state. If the successor state is newly generated, it is added to a successor state list. The generation function is called again, with the depth value decremented and the successor state list passed as the new list of available states. Execution ceases when the depth value reaches zero or the list of available states is empty.

3.3.5   Output

In order to aid understanding the generation process, this section expands on the illustrative example introduced in the preceding sections and demonstrates a sample attack graph generation process using the example. The reader is warned to avoid searching for meaning or design in the selection of these assets, qualities, and topologies; they are intended to be illustrative only, and any relation to real world networks, problems, attacks, or situations is purely coincidental. This section will use the following network model specification.

```
network model =
    assets :
```

26

```
        asset_1 ;
        asset_2 ;
        asset_3 ;


    facts :
            quality : asset_1 , quality_1=value_1 ;
            quality : asset_2 , quality_1=value_2 ;
            quality : asset_3 , quality_1=value_2 ;
            topology : asset_1 ->asset_2 , topology_1 ;
            topology : asset_2 <->asset_3 , topology_2 ;
.
```

The following exploit pattern specifications are used in this section.

```
exploit  exploit_1 ( asset_param_1 , asset_param_2)=
    preconditions :
            quality : asset_param_1 , quality_1=value_1 ;
            topology : asset_param_1 ->asset_param_2 , topology_1 ;
    postconditions :
            delete  topology : asset_param_1 ->asset_param_2 , topology_1 ;
            insert  quality : asset_param_1 , quality_1=value_2 ;
.


exploit  exploit_2 ( asset_param_1 , asset_param_2)=
    preconditions :
            quality : asset_param_1 , quality_1=value_2 ;
            quality : asset_param_2 , quality_1=value_2 ;
    postconditions :
            insert  topology : asset_param_1 <->asset_param_2 , topology_2 ;
```

Figure 3.1: State 1 (initial network state) of the illustrative discrete example

.

Fig. 3.1 represents the initial network state (denoted State 1) specified in this example. Note that Fig. 3.1 is *not* an attack graph, merely a convenient graph based representation of the example network in use here. Since there is only one quality in use in the entire example, its name is not included in this graph representation; instead, each asset is labeled with its identifier and the value of `quality_1` with the prefixing `value_` removed for brevity. Likewise, the edges that represent topologies are labeled with the topology name they represent with the prefixing `topology_` removed.

Execution of the generation process begins by specifying a maximum "depth" of generation, which will be 3 for the purposes of this exercise (although, as this attack graph converges with maximum shortest path of 3 from the initial state, this limitation is unnecessary), and by creating an initial list of states for analysis, which contains only State 1.

Generation of the next set of states begins by creating a list of all valid attacks on State 1; that is, selecting all valid bindings of assets given State 1's fact base to exploit pattern parameters given their preconditions. Three such bindings are possible from State 1: `exploit_1(asset_1, asset_2)`; `exploit2(asset_2, asset_3)`; and `exploit2(asset_2, asset_3)`. The latter two insert a topology into the fact base that already exists, therefore generating State 1, so State 1 has itself as a successor state twice. As State 1 already exists, these edges are added, and no further processing due to those transitions is done. The first attack, `exploit_1(asset_1,`

Figure 3.2: State 2 of the illustrative discrete example

`asset_2`), results in a new state, designated State 2. It is generated by performing the operations on State 1's fact base and creating a new state based on the results. State 2 is the only state added to the list of successor states.

For the next execution of the generation function, the successor state list becomes the new analysis state list, consisting now only of State 2, and the remaining allowed "depth" is decremented to 2. 6 attacks are possible from State 2: `exploit_2(asset_1, asset_2)`, `exploit_2(asset_2, asset_1)`, `exploit_2(asset_1, asset_3)`, `exploit_2(asset_3, asset_1)`, `exploit_2(asset_2, asset_3)`, and `exploit_2(asset_3, asset_2)`. The process iterates through these attacks, generating new states and, if they are new, adding them to the list of successor states. `exploit_2(asset_1, asset_2)` results in a new state, State 3 (see Fig. 3.3), which is added to the list of successor states. `exploit_2(asset_2, asset_1)` also results in State 3, which already exists and is therefore not added to the list of successor states. `exploit_2(asset_1, asset_3)` generates a new state, State 4 (see Fig. 3.4), which is added to the list of successor states. `exploit_2(asset_3, asset_1)` also generates State 4. `exploit_2(asset_2, asset_3)` and `exploit_2(asset_3, asset_2)` both result in State 2, which exists and is not added to the successor states.

The next iteration finds the remaining allowed depth at 1 and the list of analysis states to contain State 3 and State 4. From State 3, the possible attacks

Figure 3.3: State 3 of the illustrative discrete example

Figure 3.4: State 4 of the illustrative discrete example

are the same as the possible attacks in State 2: `exploit_2(asset_1, asset_2)`, `exploit_2(asset_2, asset_1)`, `exploit_2(asset_1, asset_3)`, `exploit_2(asset_3, asset_1)`, `exploit_2(asset_2, asset_3)`, and `exploit_2(asset_3, asset_2)`. Only the middle two, `exploit_2(asset_1, asset_3)`, `exploit_2(asset_3, asset_1)`, generate a state that is not State 3 itself: State 5 (see Fig. 3.5), which is added to the list of successor states. The rest simply create new reflexive edges on State 3. The list of possible attacks on State 4 is the same, with only `exploit_2(asset_1, asset_2)` and `exploit_2(asset_2, asset_1)` generating a state other than State 4 itself. Both of those attacks generate State 5 again, which already exists and therefore is not added to the list of successor states.

The next iteration finds the list of analysis states to contain only State 5, and the remaining allowed depth is 0. Because the depth limit is reached, generation ceases. However, even if the depth limit had not been reached, execution would have soon terminated. All 6 possible attacks from State 5 would have generated State 5 itself. Therefore the list of successor states would be empty, and the next iteration's list of analysis states would be empty, which would also cause generation to cease. A representation of the resulting attack graph (with execution halting after 3 iterations) is provided in Fig. 3.6.

30

Figure 3.5: State 5 of the illustrative discrete example

Figure 3.6: The illustrative example's attack graph to depth 3

3.4   Working Lexicon

3.4.1   Introduction

Because of the unrestricted nature of the terms available to a modeler in this version of the attack graph framework, in order to proceed systematically some conventions must be established in the use of terms. This thesis recommends a set of conventions designed to ease the goal of automated exploit pattern extraction from the National Vulnerability Database (NVD) maintained by the National Institute of Standards and Technology (NIST), which among other roles indexes MITRE's Common Vulnerability Enumeration (CVE). "NVD includes databases of security checklists, security related software flaws, misconfigurations, product names, and impact metrics" [2].

3.4.2   National Vulnerability Database

Several concepts are used across the NVD's index of vulnerabilities; they strongly inform the working lexicon employed in this thesis and include a vulnerability's access vector and impact type.

Access vector: The access vector of a vulnerability on NVD refers to the logical location from which the attacker may launch the attack. The possible attack vectors according to the National Vulnerability Database are as follows.

Local The vulnerability is exploitable through physical or local account access to the device on which the vulnerability resides.

Adjacent Network The vulnerability is exploitable through access to a network that is adjacent to the vulnerable host; that is, the attacker must be in the same broadcast domain or collision domain (e.g. the same network segment or VLAN).

Network No local or adjacent access is required to exploit the vulnerability; in other words, it is exploitable over the Internet.

Impact Type: The vulnerability's impact type on NVD places the effects of the vulnerability's exploitation on the target into one of the following categories. , and another for some other type of privileged access.

Confidentiality Confidentiality impacts allow the unauthorized disclosure of information (corresponding to STRIDE's "information disclosure")

Integrity Allows modification of data (corresponding to STRIDE's "tampering")

Availability Availability impacts correspond to disruptions of service (STRIDE's "denial of service")

Security Protection The security protection category refers to the effects of exploits that provide unauthorized access to the target. This may be either general system access or application access. The security protection category corresponds roughly to STRIDE's "elevation of privileges" and also partly encompasses "spoofing identity". It has three subcategories:

User access This subcategory refers to the attacker's gaining user level access to the operating system.

Administrative access This subcategory refers to the attacker's gaining root level access to the operating system.

Other access This subcategory refers to any other type of privileged access on the target.

3.4.3 Topologies

Introduction of actual terms begins with topologies, but first a note about the modeling of adversaries is needed.

Two approaches are possible for modeling the attacker. The first is to design the network model so that it is, in a way, from the attacker's perspective. In this scheme, henceforth the *first person* strategy, the attacker is treated as implicit, and its access and connections are considered properties of the system itself. The adversary's access level to a server, for instance, would be modeled as a quality of that server asset.

In the second approach, henceforth the *third person* strategy, the attacker is modeled as a first class part of the system: as an asset. Its properties, connections, and access levels are modeled as qualities and topologies of the adversary asset. This strategy, which is the one employed primarily in this thesis, can model multiple adversaries. Aside from that difference, they are roughly equivalent and a matter of taste on behalf of the modeler and analyst. This section's terminology, however, assumes the third person model.

Two types of topology terms are introduced in this section: connection and access. These correspond directly to the two NVD concepts introduced in the previous section.

Connection Topologies: Connection topologies refer to how two assets are connected, over the network or otherwise. They may be local, adjacent, or network connected. These connection topologies begin with the word `connected`, followed by an underscore, then one of the three connection types: `local`, `adjacent`, or `network`. For local and adjacent connections, this is all that is necessary.

For network connections it is necessary to specify the available protocols individually. These are done by appending another underscore, then the lower case version of the standard abbreviation for the protocol (e.g. `connected_network_http` for web or `connected_network_ssh` for secure shell).

Connections are one-way: the source of the connection topology is

Access Topologies:

3.4.4   *Qualities*

3.5   Examples

3.5.1   *Blunderdome*

3.5.2   *Denial of Sleep*

3.6   State Predicates

3.7   State Aggregation

CHAPTER 4

HYBRID EXTENSIONS

4.1 Introduction

4.2 Definition of New Syntax

4.3 Time

4.4 Time State Aggregation

CHAPTER 5

RESULTS

CHAPTER 6
CONCLUSION AND FUTURE WORK

BIBLIOGRAPHY

[1] Report: Cyber-physical systems summit. Technical report, National Science Foundation, 2008.

[2] National vulnerability database version 2.2, 2011. Web page. Retrieved April 13, 2011.

[3] R. Alur, C. Courcoubetis, T. Henzinger, and P. Ho. Hybrid automata: An algorithmic approach to the specification and verification of hybrid systems. *Hybrid systems*, pages 209–229, 1993.

[4] P. Ammann, D. Wijesekera, and S. Kaushik. Scalable, graph-based network vulnerability analysis. In *Proceedings of the 9th ACM Conference on Computer and Communications Security*, pages 217–224. ACM, 2002.

[5] M. Brownfield, Y. Gupta, and N. Davis. Wireless sensor network denial of sleep attack. In *Information Assurance Workshop, 2005. IAW'05. Proceedings from the Sixth Annual IEEE SMC*, pages 356–364. IEEE, 2005.

[6] T.K. Buennemeyer, G.A. Jacoby, W.G. Chiang, R.C. Marchany, and J.G. Tront. Battery-sensing intrusion protection system. In *Information Assurance Workshop, 2006 IEEE*, pages 176–183. IEEE.

[7] C. Campbell, J. Dawkins, B. Pollet, K. Fitch, J. Hale, and M. Papa. On Modeling Computer Networks for Vulnerability Analysis. *DBSec*, pages 233–244, 2002.

[8] K. Chen, H. Tsai, Y. Liu, and J. Shuler. A Radiofrequency Identification (RFID) Temperature-Monitoring System for Extended Maintenance of Nuclear

Materials Packaging. In *Proceedings of 2009 ASME Pressure Vessels and Piping Division Conference, Prague, Czech Republic*, 2009.

[9] T.L. Crenshaw and S. Beyer. UPBOT: a testbed for cyber-physical systems. In *Proceedings of the 3rd international conference on Cyber security experimentation and test*, pages 1–8. USENIX Association, 2010.

[10] M. Dacier and Y. Deswarte. Privilege graph: an extension to the typed access matrix model. *Computer Security ESORICS 94*, pages 319–334, 1994.

[11] G. Frehse. PHAVer: Algorithmic verification of hybrid systems past HyTech. *Hybrid Systems: Computation and Control*, pages 258–273, 2005.

[12] T.A. Henzinger. The theory of hybrid automata. In *Logic in Computer Science, 1996. LICS'96. Proceedings., Eleventh Annual IEEE Symposium on*, pages 278–292. IEEE, 1996.

[13] T.A. Henzinger, P.H. Ho, and H. Wong-Toi. HyTech: A model checker for hybrid systems. *International Journal on Software Tools for Technology Transfer (STTT)*, 1(1):110–122, 1997.

[14] T.A. Henzinger, P.W. Kopke, A. Puri, and P. Varaiya. What's decidable about hybrid automata? *Journal of Computer and System Sciences*, 57(1):94–124, 1998.

[15] K. Ingols, M. Chu, R. Lippmann, S. Webster, and S. Boyer. Modeling modern network attacks and countermeasures using attack graphs. In *2009 Annual Computer Security Applications Conference*, pages 117–126. IEEE, 2009.

[16] D.K. Kaynar, N. Lynch, R. Segala, and F. Vaandrager. The theory of timed I/O automata. *Synthesis Lectures on Distributed Computing Theory*, 1(1):1–137, 2010.

[17] E.A. Lee. Cyber-physical systems-are computing foundations adequate. In *Position Paper for NSF Workshop On Cyber-Physical Systems: Research Motivation, Techniques and Roadmap*. Citeseer, 2006.

[18] R.P. Lippmann, K.W. Ingols, and MASSACHUSETTS INST OF TECH LEXINGTON LINCOLN LAB. *An annotated review of past papers on attack graphs*. Massachusetts Institute of Technology, Lincoln Laboratory, 2005.

[19] G. Louthan, W. Roberts, M. Butler, and J. Hale. The blunderdome: an offensive exercise for building network, systems, and web security awareness. In *Proceedings of the 3rd international conference on Cyber security experimentation and test*, pages 1–7. USENIX Association, 2010.

[20] J. Lygeros, K.H. Johansson, S. Sastry, and M. Egerstedt. On the existence of executions of hybrid automata. In *Decision and Control, 1999. Proceedings of the 38th IEEE Conference on*, volume 3, pages 2249–2254. IEEE, 1999.

[21] N. Lynch, R. Segala, and F. Vaandrager. Hybrid I/O automata revisited. *Hybrid Systems: Computation and Control*, pages 403–417, 2001.

[22] N. Lynch, R. Segala, F. Vaandrager, and H. Weinberg. Hybrid I/O automata. *Hybrid Systems III*, pages 496–510, 1996.

[23] C. Phillips and L.P. Swiler. A graph-based system for network-vulnerability analysis. In *Proceedings of the 1998 workshop on New security paradigms*, pages 71–79. ACM, 1998.

[24] D.R. Raymond, R.C. Marchany, M.I. Brownfield, and S.F. Midkiff. Effects of denial-of-sleep attacks on wireless sensor network MAC protocols. *Vehicular Technology, IEEE Transactions on*, 58(1):367–380, 2009.

[25] B. Schneier. Modeling security threats. *Dr. Dobb's journal*, 24(12), 1999.

[26] O. Sheyner, J. Haines, S. Jha, R. Lippmann, and J.M. Wing. Automated generation and analysis of attack graphs. 2002.

[27] S.J. Templeton and K. Levitt. A requires/provides model for computer attacks. In *Proceedings of the 2000 workshop on New security paradigms*, pages 31–38. ACM, 2001.

[28] T. Tidwell, R. Larson, K. Fitch, and J. Hale. Modeling internet attacks. In *Proceedings of the 2001 IEEE Workshop on Information Assurance and security*, volume 59, 2001.

[29] L. Wang, S. Noel, and S. Jajodia. Minimum-cost network hardening using attack graphs. *Computer Communications*, 29(18):3812–3824, 2006.